

Deployment
o

Server configuration
oooooooooooooooooooo
oooooooooooo

Security
o
o
oooooo

Extras
o

8 Deployment

CS 425

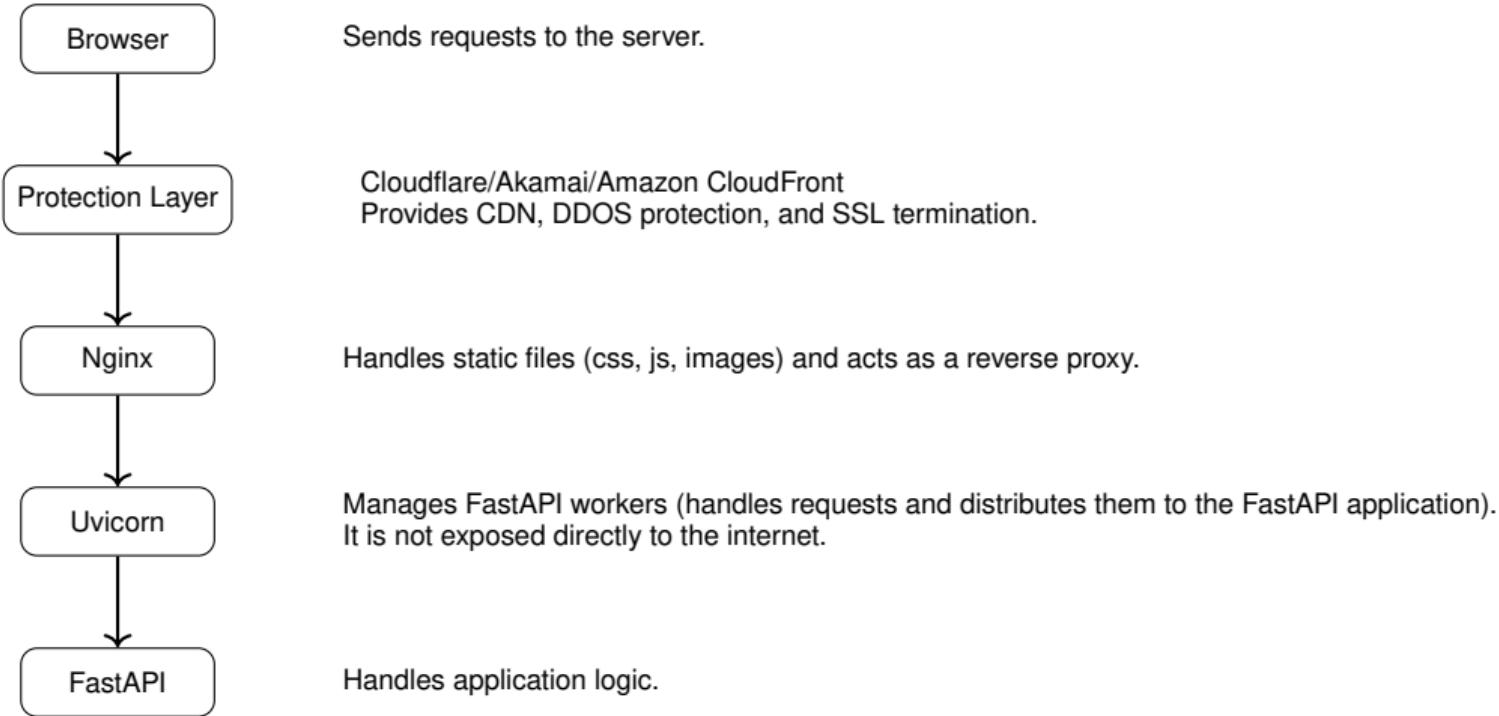
Web Applications Development

Alex Petrovici

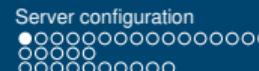
Content

- 1 Deployment
 - 2 Server configuration
 - Create a new server
 - Setup the server
 - Setup unicorn and nginx
 - 3 Security
 - SSH security
 - Firewall configuration
 - Control access
 - 4 Extras

Deployment / Introduction



[Create a new server](#)

Server configuration
A progress bar consisting of a single black dot followed by nine white circles.

A progress bar consisting of five white circles.

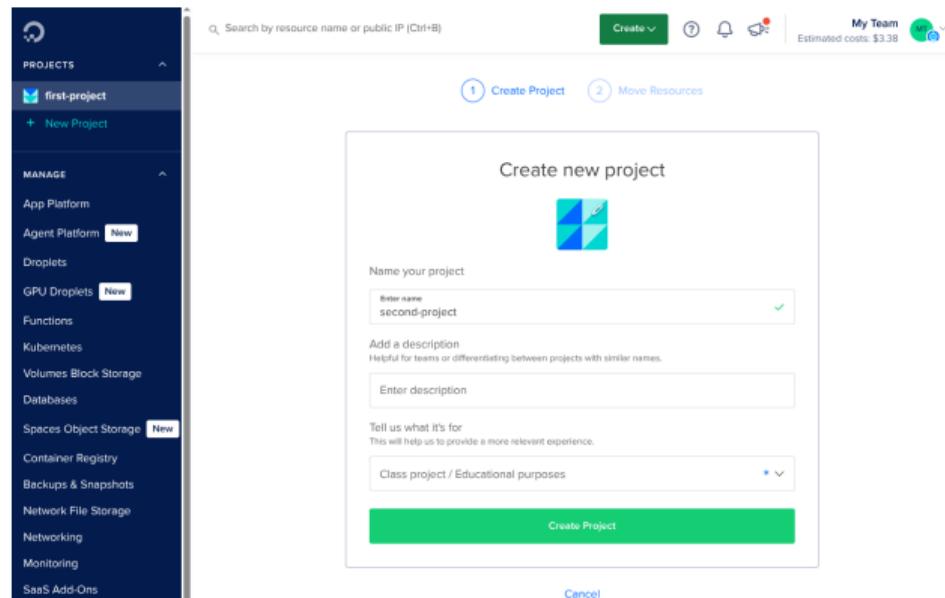
Server configuration / Create a new server

- ① Create a new server
- ② Setup users
- ③ Setup the server
- ④ Setup unicorn
- ⑤ Setup nginx

Create a new server

Server configuration / Create a new server

- ① Make an account in DigitalOcean.
 - ② Make a new project.



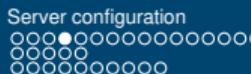
Create a new server

Server configuration / Create a new server

- ③ Make the server. (Go to Create / Droplets)

The screenshot shows the DigitalOcean dashboard with the following elements:

- Header:** "Search by resource name or public IP (Ctrl+B)" and "Create" button.
- Project Overview:** "second-project" (Class project / Educational purposes) with a thumbnail icon.
- Navigation:** "Resources" (selected), "Activity", "Settings".
- Welcome Section:** "Welcome to DigitalOcean!" message, "Getting Started" (with a book icon), and "Spin up a Droplet" (with a droplet icon).
- Resource Catalog:** A sidebar with various service options:
 - GPU Droplets**: Create cloud servers with GPUs.
 - Droplets**: Create cloud servers.
 - Agents**: Create AI Agents.
 - Serverless Inference**: Access AI Models.
 - Knowledge Bases**: Create AI Knowledge Bases.
 - Kubernetes**: Create Kubernetes clusters.
 - App Platform**: Deploy your code.
 - Functions**: Create Cloud Functions.
 - SaaS Add-Ons**: Deploy Marketplace software.
 - Databases**: Create database clusters.
 - Volumes Block Storage**: Add storage to Droplets.
- Team & Costs:** "My Team" (1 member), "Estimated costs: \$3.38", and a "Move Resources" button.
- Cloudways & DigitalOcean:** A callout box stating "#1 Managed Hosting Provider joins forces with DigitalOcean, check it out!" with icons for Cloudways and DigitalOcean.



Create a new server

Server configuration / Create a new server

- ④ Choose the region closest to your users.

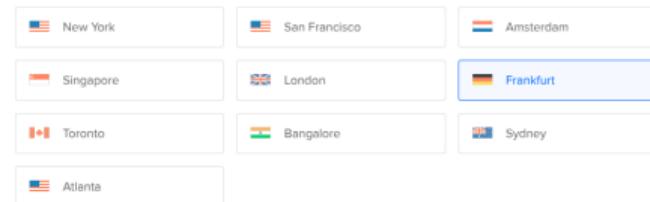


Create Droplets

Learn

Droplets are virtual machines that anyone can setup in seconds. You can use droplets, either standalone or as part of a larger, cloud based infrastructure.

Choose Region



Datacenter

Frankfurt • Datacenter 1 • FRA1

◆ Tip: Select the datacenter closest to you or your users

Dismiss

Avoid any potential latency by selecting a region closest to you - a region is a geographic area where we have one or more datacenters.

[Create a new server](#)

Server configuration / Create a new server

- ⑤ Choose the OS you prefer (recommended Ubuntu or Debian).

Choose an image

OS Marketplace (263) Custom images



Version

13 x64



Create a new server

Server configuration / Create a new server

- ⑥ Choose the size of the server. (We can start with the smallest one. For reference, Debian can run with at least 512MB of RAM, but is recommended to use at least 1GB of RAM.)

Choose Size Need help picking a plan? [Help me choose](#)

Droplet Type

DEDICATED CPU			
General Purpose	CPU-Optimized	Memory-Optimized	Storage-Optimized
<input checked="" type="radio"/> Basic (Plan selected)	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>

Basic virtual machines with a mix of memory and compute resources. Best for small projects that can handle variable levels of CPU performance, like blogs, web apps and dev/test environments.

CPU options

<input checked="" type="radio"/> Regular Disk type: SSD	<input type="radio"/> Premium Intel Disk: NVMe SSD	<input type="radio"/> Premium AMD Disk: NVMe SSD
---	--	--

\$ 4/mo \$0.006/hour	\$ 6/mo \$0.009/hour	\$ 12/mo \$0.018/hour	\$ 18/mo \$0.027/hour	\$ 24/mo \$0.036/hour	\$ 48/mo \$0.071/hour
512 MB / 1 CPU 10 GB SSD Disk 500 GB transfer	1 GB / 1 CPU 25 GB SSD Disk 1000 GB transfer	2 GB / 1 CPU 50 GB SSD Disk 2 TB transfer	2 GB / 2 CPUs 60 GB SSD Disk 3 TB transfer	4 GB / 2 CPUs 80 GB SSD Disk 4 TB transfer	8 GB / 4 CPUs 160 GB SSD Disk 5 TB transfer

Create a new server

Server configuration / Create a new server

- ⑦ Now make a ssh key to access the server.

```
> cd ~/.ssh
> ssh-keygen
Generating public/private ed25519 key pair.
Enter file in which to save the key (/home/alex/.ssh/id_ed25519): class_project
Enter passphrase for "class_project" (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in class_project
Your public key has been saved in class_project.pub
The key fingerprint is:
SHA256:kaRDw05lstihC204zbIvUSRfuKx5kLr1E7K420U1WHY alex@zenbook-debian
The keyls randomart image is:
+--[ED25519 256]--+
|   o... .
| . o o. o .
| * +  o o
|*.ooo. .
|o+o .. |
+---[SHA256]---+
> cat class_project.pub | wl-copy
```

Create a new server

Server configuration
○○○○○○●○○○○○○○
○○○○○○○○○○

Security
○
○○○○○

Server configuration / Create a new server

- ⑧ Copy the public key to the server.

Choose Authentication Method 

SSH Key
Connect to your Droplet with an SSH key pair

Password
Connect to your Droplet as the "root" user via password

Choose your SSH keys  Select at least one key.

Select all desktop digital-cl... digital-oc... stc-laptop...

New SSH Key

Add public SSH key

Copy your public SSH key and paste it in the space below. For instructions on how, follow the steps on the right.

SSH key content
ssh-ed25519
AAAAC3NzaC1EzDHNTE5AAAAInQ/JKpJ0MnRFCGybINjLyLewyYYc7HDQpB
mrU7QRlJ alex@zenbook-debian

Name
class_project

Add SSH Key

SSH Keys

Follow these instructions to create or add SSH keys on Linux, MacOS & Windows.
Windows users without OpenSSH can install and use PuTTY instead.

Create a new key pair, if needed

Open a terminal and run the following command:

ssh-keygen 

You will be prompted to save and name the key:

Generating public/private rsa key pair. Enter file in which to save the key

[Create a new server](#)

Server configuration / Create a new server

- ⑨ Finnish creating the server.

Finalize Details

Quantity

Deploy multiple Droplets with the same configuration.

Hostname

Give your Droplets an identifying name you will remember them by.

Tags

Project

\$6.00/month

\$0.009/hour

[CREATE VIA COMMAND LINE](#)[Create Droplet](#)

Create a new server

Server configuration / Create a new server

- ⑩ Wait until the server is ready.



Server configuration / Create a new server

Setup users

- ① Enter in the server as root.

The end goal is to connect via ssh as a user (**NOT** as a root).

To do so, first we need to create this new user. Initially we will enter as root, make the new user, and from there only allow to enter as the user we just created.

Use the previous ip address and the **private** key we created earlier.

```
> ssh root@207.154.253.161 -i ~/.ssh/class_project
Linux debian-s-1vcpu-1gb-fra1-01 6.12.43+deb13-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.12.43-1
→ (2025-08-27) x86_64
```

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms **for** each program are described **in** the
individual files **in** /usr/share/doc/*/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

[Create a new server](#)

Server configuration / Create a new server

- ② Add the new user.

```
adduser <username>
```

In practice, add the user as follows:

```
root@debian-s-1vcpu-1gb-fra1-01:~# adduser project
New password:
Retype new password:
passwd: password updated successfully
Changing the user information for project
Enter the new value, or press ENTER for the default
  Full Name []: project
  Room Number []:
  Work Phone []:
  Home Phone []:
  Other []:
Is the information correct? [Y/n] y
root@debian-s-1vcpu-1gb-fra1-01:~#
```

Server configuration / Create a new server

- ③ Add the user to the sudoers file so we can use sudo and install packages.

```
usermod -aG sudo <username>
```

- **-a:** **append** the user to the supplementary group(s) specified by the **-G** flag. Without this, the user would be removed from all other supplementary groups not listed.
- **-G sudo:** Add the user to the sudo group.
- **<username>:** The user to be modified.

In summary, this command adds the user project to the sudo group, allowing them to execute commands as superuser.

In practice, add the user as follows:

```
root@debian-s-1vcpu-1gb-fra1-01:~# usermod -aG sudo project
```

Server configuration / Create a new server

- ④ Change the ssh configuration to allow only the new user to login via ssh key.

Since we made the server from digitalocean, we already have the public key loaded for the root user. You can check it by running:

```
cat /root/.ssh/authorized_keys
```

Now we will allow the new user <username> to login by using this ssh key.

```
mkdir -p /home/<username>/.ssh
cp ~/.ssh/authorized_keys /home/<username>/.ssh/
chown -R <username>:<username> /home/<username>/.ssh
chmod 700 /home/<username>/.ssh
chmod 600 /home/<username>/.ssh/authorized_keys
```

- `mkdir -p /home/<username>/.ssh`: Creates the user's .ssh directory if it doesn't exist.
- `cp ~/.ssh/authorized_keys /home/<username>/.ssh/`: Copies your current user's public key(s) to the new user's authorized keys file, allowing SSH login with your key.
- `chown -R <username>:<username> /home/<username>/.ssh`: Changes ownership of the .ssh directory and its contents to the new user, ensuring they have the correct permissions.
- `chmod 700 /home/<username>/.ssh`: Sets the directory permissions so only the user can read, write, execute.
- `chmod 600 /home/<username>/.ssh/authorized_keys`: Sets the authorized keys file to be readable and writable only by the user for security.

[Create a new server](#)

Server configuration / Create a new server

So, in practice, we will run the following commands:

```
root@debian-s-1vcpu-1gb-fra1-01:~# mkdir -p /home/project/.ssh
root@debian-s-1vcpu-1gb-fra1-01:~# cp ~/.ssh/authorized_keys /home/project/.ssh/
root@debian-s-1vcpu-1gb-fra1-01:~# chown -R project:project /home/project/.ssh
root@debian-s-1vcpu-1gb-fra1-01:~# chmod 700 /home/project/.ssh
root@debian-s-1vcpu-1gb-fra1-01:~# chmod 600 /home/project/.ssh/authorized_keys
```

Server configuration / Create a new server

- ⑤ Check that you can login with the new user.

```
systemctl restart ssh
```

Now, in the client, in a separate terminal, try to log in with the new user.

```
> ssh project@207.154.253.161 -i ~/.ssh/class_project
Linux debian-s-1vcpu-1gb-fra1-01 6.12.43+deb13-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.12.43-1
→ (2025-08-27) x86_64
```

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms **for** each program are described **in** the
individual files **in** /usr/share/doc/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
project@debian-s-1vcpu-1gb-fra1-01:~\$

Since we connected successfully as <username>, then exit the root session.

Setup the server

Server configuration / Setup the server

- ① Update the system.

```
sudo apt update && sudo apt upgrade -y
```

- ② Give some swap space to the server. Read more about the process at <https://www.digitalocean.com/community/tutorials/how-to-add-swap-space-on-debian-11>

```
sudo fallocate -l 2G /swapfile
sudo chmod 600 /swapfile
sudo mkswap /swapfile
sudo swapon /swapfile
sudo swapon --show
```

Then make permanent by adding the file to the fstab file.

```
sudo cp /etc/fstab /etc/fstab.bak
echo '/swapfile none swap sw 0 0' | sudo tee -a /etc/fstab
```

Setup the server

Server configuration / Setup the server

- `fallocate -l 2G /swapfile`: Allocates a file of exactly 2 GiB by reserving blocks in the filesystem metadata. Fast. Does not write zeros; it just marks space as used.
- `chmod 600 /swapfile`: Make the file only accessible to root only.
- `mkswap /swapfile`: Writes swap metadata (header, UUID, page map). Without this, the kernel does not recognize the file as swap.
- `swapon /swapfile`: Registers the file with the kernel and starts using it as swap.
- `swapon --show`: Verification. Shows active swap devices/files.

In practice, we will run the following commands:

```
root@debian-s-1vcpu-1gb-fra1-01:~# sudo fallocate -l 2G /swapfile
root@debian-s-1vcpu-1gb-fra1-01:~# sudo chmod 600 /swapfile
root@debian-s-1vcpu-1gb-fra1-01:~# sudo mkswap /swapfile
Setting up swapspace version 1, size = 2 GiB (2147479552 bytes)
no label, UUID=090997d6-3fd9-4e7c-876d-8abba6289ad0
root@debian-s-1vcpu-1gb-fra1-01:~# sudo swapon /swapfile
root@debian-s-1vcpu-1gb-fra1-01:~# sudo swapon --show
NAME      TYPE SIZE USED PRIO
/swapfile file  2G   0B   -2
```

Setup the server

Server configuration / Setup the server

- ④ Install required system packages:

```
sudo apt install python3 python3-pip python3-venv nginx git pipx ncurses-term nano -y
pipx ensurepath
pipx install poetry
pipx ensurepath
```

Then close the terminal and connect again.

- ⑤ Inside the server, link the server to your github account using ssh (using ssh-keygen).
- ⑥ Then clone your project repository.

```
git clone <repository-url>
git clone git@github.com:agpetrovici/fastapi_project.git
cd fastapi_project
```

- ⑦ Install the dependencies.

```
python3 -m venv ./venv
source ./venv/bin/activate
poetry install --only main --no-interaction --no-ansi
```

Setup the server

Server configuration / Setup the server

Sanity check: test whether the server can run the application.

```
source ./venv/bin/activate
python3 app/backend/main.py
```

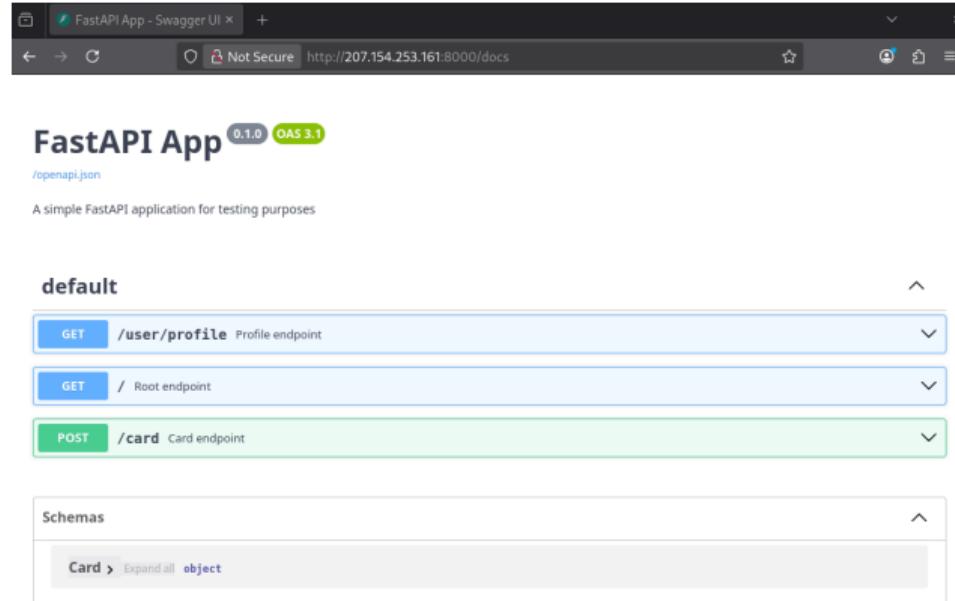
Expected output:

```
(.venv) project@debian-s-1vcpu-1gb-fra1-01:~/fastapi_project$ python3 app/backend/main.py
INFO:     Started server process [11842]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
INFO:     86.127.228.177:55316 - "GET / HTTP/1.1" 307 Temporary Redirect
INFO:     86.127.228.177:55316 - "GET /docs HTTP/1.1" 200 OK
INFO:     86.127.228.177:55316 - "GET /openapi.json HTTP/1.1" 200 OK
```

Setup the server

Server configuration / Setup the server

In a browser, go to `http://<server-ip>:8000` and you should see it working:



This would work as long as the terminal is working. Next step is to configure a systemd

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

Setup unicorn

We will setup nginx and unicorn to run the application automatically.

First, check that unicorn works.

```
uvicorn app.backend.application:create_app --factory --reload --host 0.0.0.0 --port 8000
```

Expected output:

```
(.venv) project@debian-s-1vcpu-1gb-fra1-01:~/fastapi_project$ pwd  
/home/project/fastapi_project  
(.venv) project@debian-s-1vcpu-1gb-fra1-01:~/fastapi_project$ uvicorn  
→ app.backend.application:create_app --factory --reload --host 0.0.0.0 --port 8000  
INFO:     Will watch for changes in these directories: ['/home/project/fastapi_project']  
INFO:     Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)  
INFO:     Started reloader process [11861] using WatchFiles  
INFO:     Started server process [11863]  
INFO:     Waiting for application startup.  
INFO:     Application startup complete.
```

Now that it works, we will make the systemd service.

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

In production is better to use `unicorn.sock` instead of specifying the host and port, so let's try that.

First, check that `unicorn.sock` works.

```
uvicorn app.backend.application:create_app --factory --uds
→ /home/project/fastapi_project/unicorn.sock
```

Expected output:

```
(.venv) project@debian-s-1vcpu-1gb-fra1-01:~/fastapi_project$ pwd
/home/project/fastapi_project
(.venv) project@debian-s-1vcpu-1gb-fra1-01:~/fastapi_project$ uvicorn
→ app.backend.application:create_app --factory --uds
→ /home/project/fastapi_project/unicorn.sock
INFO:     Started server process [12305]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on unix socket /home/project/fastapi_project/unicorn.sock (Press
→ CTRL+C to quit)
```

Now that it works, we will make the systemd service.

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

```
sudo nano /etc/systemd/system/project.service
```

Add the following service description:

```
[Unit]
Description=Uvicorn instance to serve the project app
After=network.target

[Service]
User=project
Group=www-data
WorkingDirectory=/home/project/fastapi_project
Environment="PATH=/home/project/fastapi_project/.venv/bin"
ExecStart=/home/project/fastapi_project/.venv/bin/uvicorn app.backend.application:create_app
→ --factory --reload --uds /home/project/fastapi_project/uvicorn.sock

[Install]
WantedBy=multi-user.target
```

Save and exit.

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

Then start the service.

```
sudo systemctl start project
```

At this point, you can't access the app yet since we need nginx to redirect traffic through the socket. But if you use the host and port option, then you should be able to access the application at <http://<server-ip>:8000>.

Then enable the service to start on boot.

```
sudo systemctl enable project
```

Server configuration / Setup unicorn and nginx

```
project@debian-s-1vcpu-1gb-fra1-01:~$ systemctl status project
● project.service - Uvicorn instance to serve the project app
   Loaded: loaded (/etc/systemd/system/project.service; disabled; preset: enabled)
   Active: active (running) since Sun 2025-12-14 18:34:03 UTC; 3min 3s ago
     Invocation: cc4da9426d1b45f6a3d2e6a541f5755
   Main PID: 12055 (uvicorn)
      Tasks: 5 (limit: 1110)
     Memory: 54.6M (peak: 56.3M)
        CPU: 1.744s
      CGroup: /system.slice/project.service
              └─12055 /home/project/fastapi_project/.venv/bin/python /home/project/fastapi_project/.venv/bin/python -c "from multiprocessing.resource_tracker import resource_tracker"
                  ├─12056 /home/project/fastapi_project/.venv/bin/python -c "from multiprocessing.resource_tracker import resource_tracker"
                  ├─12057 /home/project/fastapi_project/.venv/bin/python -c "from multiprocessing.spawn import spawn"

Dec 14 18:34:03 debian-s-1vcpu-1gb-fra1-01 uvicorn[12055]: INFO: Will watch for changes in these directo...
Dec 14 18:34:03 debian-s-1vcpu-1gb-fra1-01 uvicorn[12055]: INFO: Uvicorn running on http://0.0.0.0:8000
Dec 14 18:34:03 debian-s-1vcpu-1gb-fra1-01 uvicorn[12055]: INFO: Started reloader process [12055] using ...
Dec 14 18:34:04 debian-s-1vcpu-1gb-fra1-01 uvicorn[12057]: INFO: Started server process [12057]
Dec 14 18:34:04 debian-s-1vcpu-1gb-fra1-01 uvicorn[12057]: INFO: Waiting for application startup.
Dec 14 18:34:04 debian-s-1vcpu-1gb-fra1-01 uvicorn[12057]: INFO: Application startup complete.
Dec 14 18:34:07 debian-s-1vcpu-1gb-fra1-01 uvicorn[12057]: INFO: 86.127.228.177:57136 - "GET /docs HTTP/1.1"
Dec 14 18:34:08 debian-s-1vcpu-1gb-fra1-01 uvicorn[12057]: INFO: 86.127.228.177:57136 - "GET /openapi.json"
lines 1-21/21 (END)
```

Now we can control the behaviour using a service, which is great because we can enter and exit the server without depending on the client terminal.

Server configuration / Setup unicorn and nginx

Fixing issues

If we make mistakes in the configuration file, then the service will not start (and we won't see a traceback since we're not running the script from the terminal).

If the service is active, (running) but it has an internal error, it won't appear in the status, but we can check the detailed logs with

```
sudo journalctl -u project
```

Then we would need to update `project.service` file, and after changing it, we need to reload `systemctl` and then restart the service.

```
sudo systemctl daemon-reload  
sudo systemctl restart project
```

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

Setup nginx

We will setup nginx to connect external traffic to the unicorn service.

```
sudo nano /etc/nginx/sites-available/project.conf
```

Add the following service description:

```
server {  
    listen 80;  
    server_name 207.154.253.161;  
    location / {  
        include proxy_params;  
        proxy_pass http://unix:/home/project/fastapi_project/uvicorn.sock;  
    }  
}
```

Save and exit.

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

We need fix the permissions of the socket file so nginx can access it.

```
sudo chmod 755 /home/project
sudo chmod 755 /home/project/fastapi_project
```

Then restart the unicorn service to recreate the socket with correct permissions.

```
sudo systemctl restart project
```

Now enable the configuration (makes a symlink to the file in the sites-enabled directory).

```
sudo ln -s /etc/nginx/sites-available/project.conf /etc/nginx/sites-enabled/
```

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

Test the configuration.

```
sudo nginx -t
```

Expected output:

```
project@debian-s-1vcpu-1gb-fra1-01:~$ sudo nginx -t
nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
nginx: configuration file /etc/nginx/nginx.conf test is successful
```

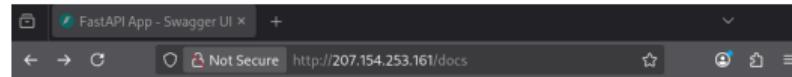
Now restart nginx to load the new configuration.

```
sudo systemctl restart nginx
```

Setup unicorn and nginx

Server configuration / Setup unicorn and nginx

Now you should be able to access the application at <http://<server-ip>>.



FastAPI App 0.1.0 OAS 3.1

</openapi.json>

A simple FastAPI application for testing purposes

default

<button>GET</button>	/user/profile	Profile endpoint	<button>▼</button>
<button>GET</button>	/	Root endpoint	<button>▼</button>
<button>POST</button>	/card	Card endpoint	<button>▼</button>

Schemas

Card > [Expand all object](#)

Security / SSH security

Disable remote login and login via password.

First edit the ssh configuration file.

```
sudo nano /etc/ssh/sshd_config
```

Add the following lines:

```
PermitRootLogin no  
PasswordAuthentication no
```

Save and exit. Then restart the ssh service.

```
sudo systemctl restart ssh
```

Security / Firewall configuration

Block unused ports.

Since this is a production server that serves a website (80, 443) and ssh (22), then we will only allow these ports to be accessible from the internet.

We will use **ufw** (uncomplicated firewall) to manage the firewall.

```
sudo apt-get install ufw -y
sudo ufw default deny incoming
sudo ufw default allow outgoing
sudo ufw allow 22/tcp
sudo ufw allow 80/tcp
sudo ufw allow 443/tcp
sudo ufw enable # This command activates the firewall with the rules you just configured.
```

Control access

Security / Control access

Ban attackers that try to log in over ssh.

```
sudo apt update  
sudo apt install fail2ban -y
```

Make a jail to control ssh access.

```
sudo nano /etc/fail2ban/jail.d/sshd.local
```

Security / Control access

Add the following lines:

```
[DEFAULT]
bantime      = -1
findtime     = 300
maxretry     = 1
backend       = systemd
loglevel     = INFO
logtarget    = /var/log/fail2ban.log

[sshd]
enabled      = true
port         = 22
filter       = sshd
logpath      = /var/log/auth.log
```

Save and exit. Then reload to put it into production `sudo fail2ban-client reload` or
`sudo systemctl restart fail2ban`

Security / Control access

Now you can check the logs to see if any attackers are being banned.

```
sudo fail2ban-client status sshd
```

Expected output:

```
project@debian-s-1vcpu-1gb-fra1-01:~$ sudo fail2ban-client status sshd
Status for the jail: sshd
|- Filter
| |- Currently failed: 0
| |- Total failed: 8
| `-' Journal matches: _SYSTEMD_UNIT=ssh.service + _COMM=sshd
`- Actions
  |- Currently banned: 1
  |- Total banned: 1
  `-' Banned IP list: 64.23.135.179
```

This shows fail2ban is already blocking unauthorized SSH login attempts.

To prevent an ever-growing permanent ban list, set bantime to a value greater than 0.

Control access

Security / Control access

If you want to block all access to the IPs detected by the sshd jail (not just SSH access), you can configure fail2ban to use firewall rules that apply system-wide.

In the `sshd.local` file, set the action to use the `iptables` banning for all ports (instead of just restricting SSH):

```
[sshd]
enabled    = true
port       = 22
filter     = sshd
logpath   = /var/log/auth.log
action     = iptables-allports[name=SSHD, protocol=all]
```

This configuration uses the `iptables-allports` action, which blocks traffic from banned IPs across all ports, not just SSH.

Security / Control access

After making this change, reload the fail2ban configuration:

```
sudo fail2ban-client reload
```

Now, any IP that gets banned by the sshd jail will be blocked from accessing *any* port on the server, not just SSH.

Control access

Security / Control access

If you are trying to test the firewall rules and you find yourself locked out of the server, you can unlock it by running the following command:

```
sudo fail2ban-client set <jail-name> unbanip <ip-address>
```

For example:

```
sudo fail2ban-client set sshd unbanip 170.39.218.202
```

NOTE: You can use <https://www.abuseipdb.com/> to check if an ip is associated with other risks.

Extras /

- Setup fail2ban to block access of ips that are trying to access sensitive files directly like /etc/passwd or /etc/shadow .
- Setup SSL so users can access the application using https (port 443).
- Setup a custom domain so users can access the application using a custom domain (e.g. <https://www.example.com>).
- Setup Cloudflare or another CDN & Security layer.